

# THML-Powered RBM Demonstration

*Training an RBM on binary bars and stripes data*

Alen Ribić & Alessio Toniolo, 2025

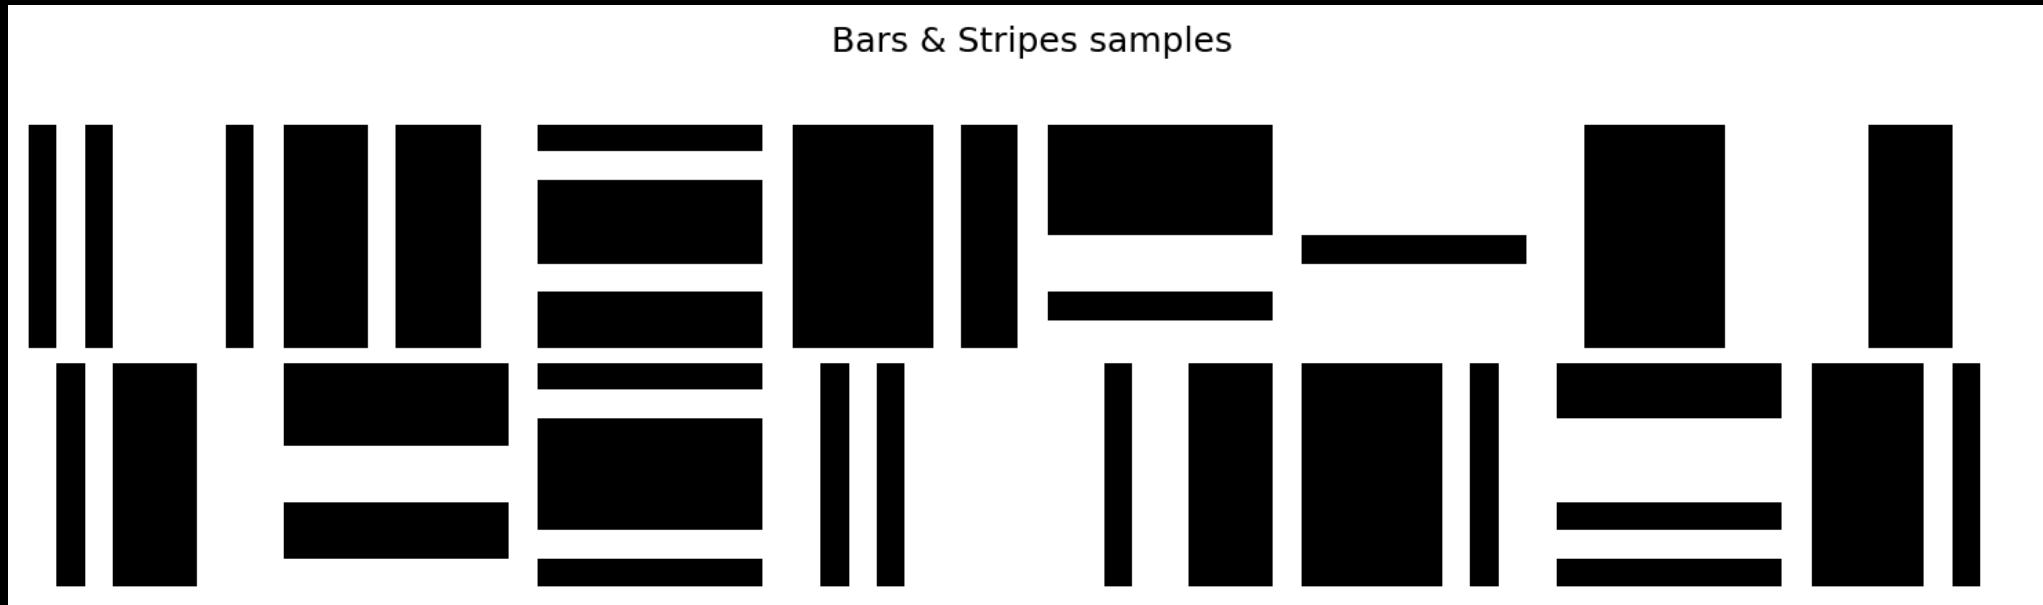
# Overview

- Build a **Restricted Boltzmann Machine (RBM)** as an **Ising Energy-Based Model**
- Train it using **THRML's KL Gradient Estimator**
- Evaluate using:
  - **1-step clamped reconstructions**
  - **THRML free-running sampling**
  - **Naive Python Gibbs sampling**
- Show how **THRML acts as a general Gibbs engine**
- Visualize learned filters and sampling results

# Dataset: 8×8 Bars & Stripes

- Synthetic binary dataset used in EBM research
- Contains:
  - Horizontal bars
  - Vertical bars
  - Multiple bars per image
- Structured, multimodal, ideal for generative testing

Bars & Stripes samples



# Training with THML

THML provides:

- **Block sampling engine**
- **KL divergence gradient estimator**
- **Factor graph representation**
- **Sampling schedules**
- Multi-chain sampling via `batch_shape`

We train RBM parameters:

- weights (  $W$  )
- visible biases (  $b_v$  ), hidden biases (  $b_h$  )
- temperature parameter (  $\beta$  )

## Our RBM/Ising energy function

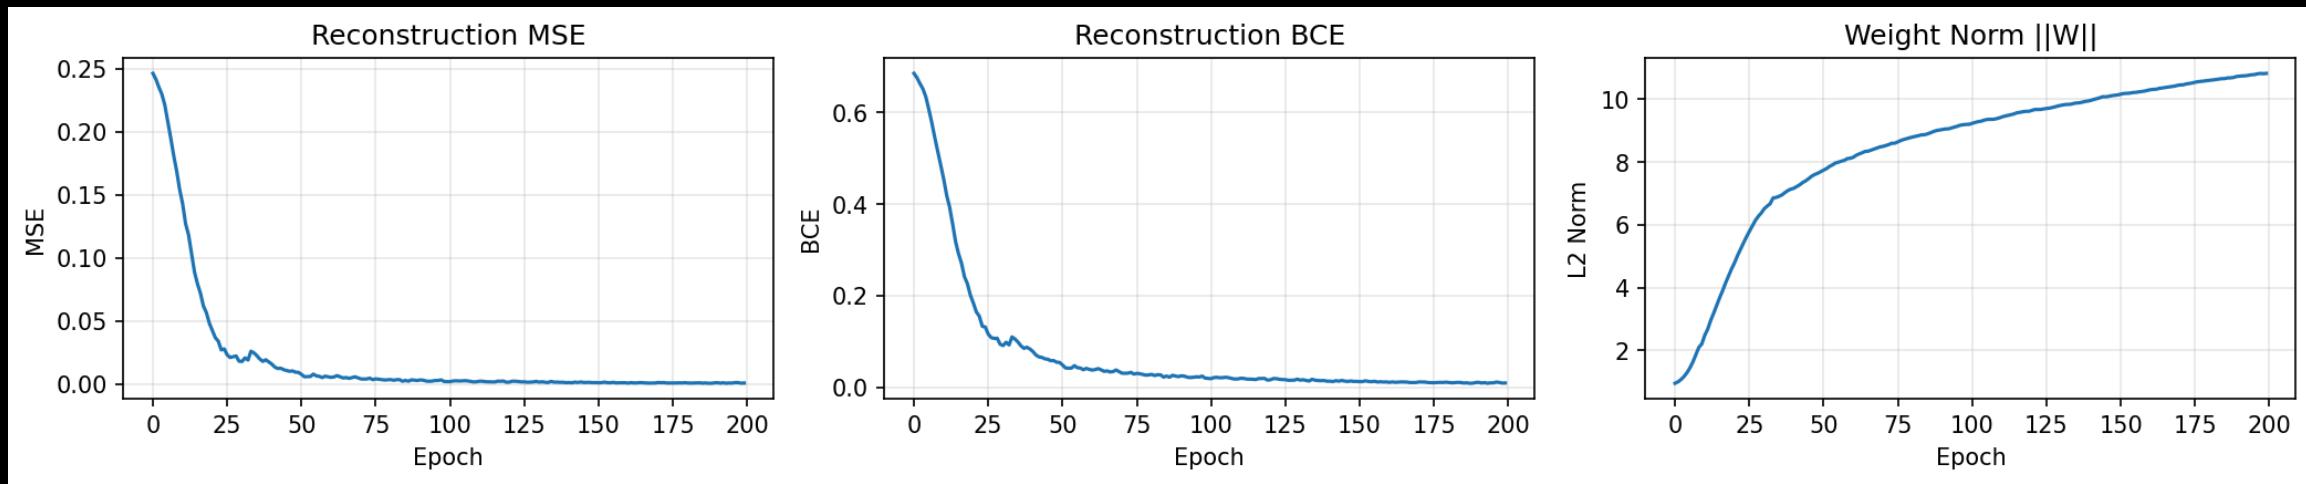
$$E(v, h) = -\beta \left[ \sum_i b_i^{(v)} v_i + \sum_j b_j^{(h)} h_j + \sum_{i,j} W_{ij} v_i h_j \right]$$

directly taken from the Gibbs/Ising Hamiltonian, and each sum corresponds to a physical energy term

- 1st term, external field on visible spins
- 2nd term, external field on hidden spins
- Pairwise interaction energy

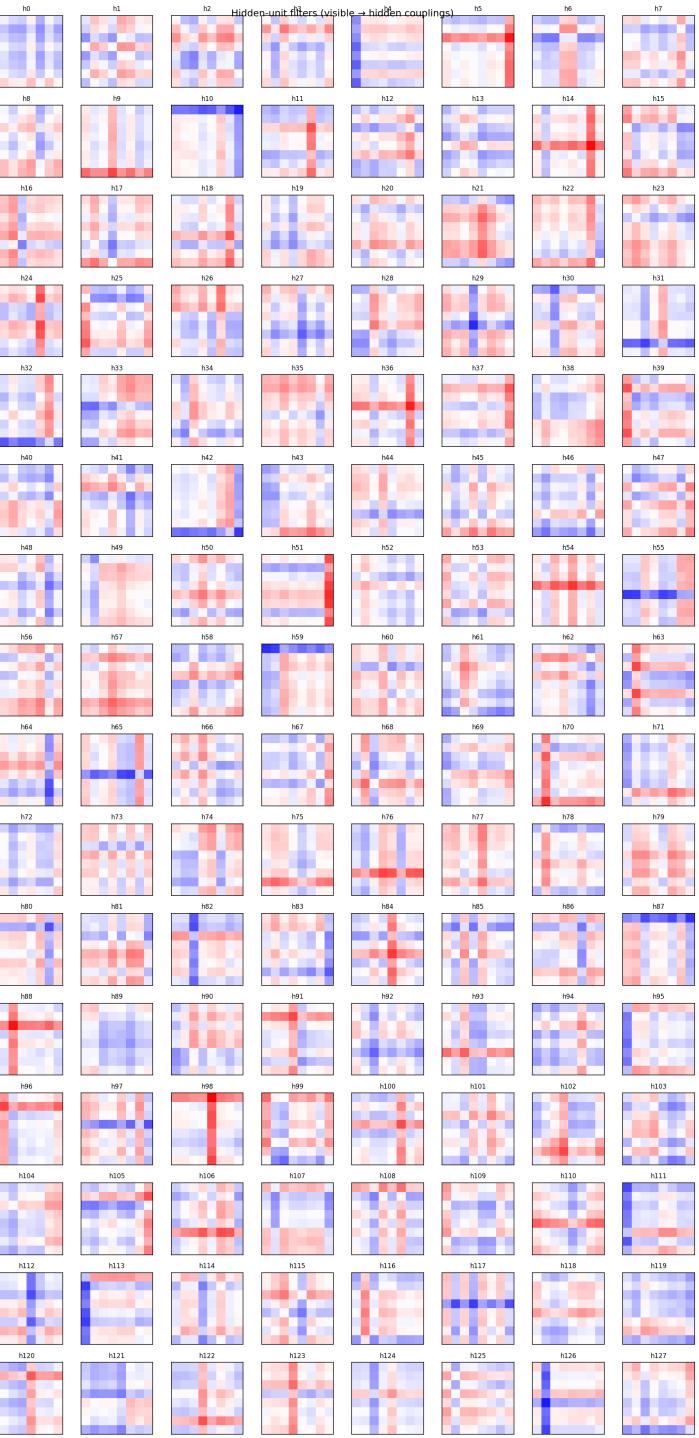
# Training Curves

- **MSE** drops sharply → good reconstructions
- **BCE** approaches ~0 → near-perfect predictions
- Weight norm increases steadily → filters specialize



# Learned Hidden Unit Filters

- Each hidden unit learns a **feature detector**
- Filters resemble:
  - Vertical bars
  - Horizontal bars
  - Crossings
  - Composite stripe patterns
- RBM learns a **distributed encoding** of the dataset



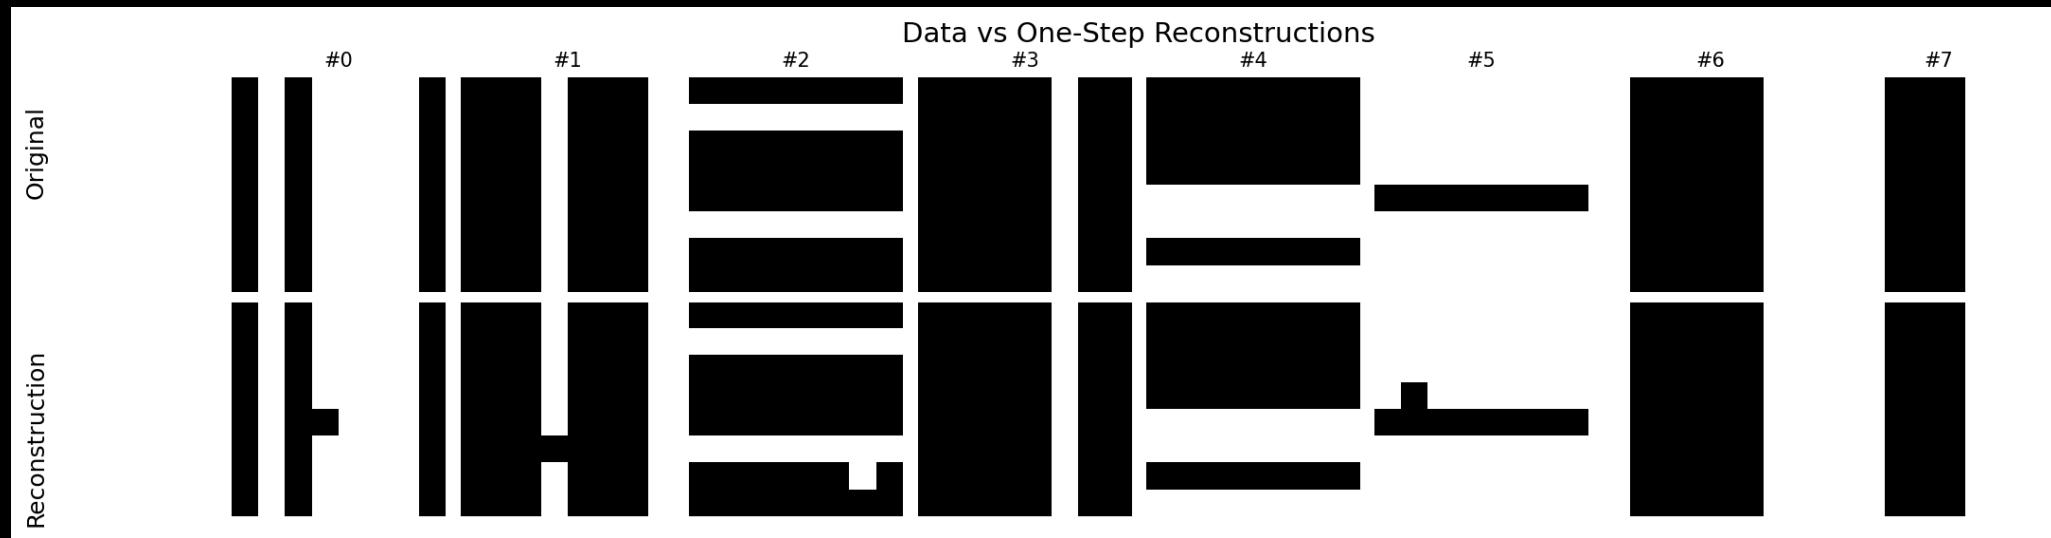
# Clamped 1-Step Reconstructions

Process:

1. Clamp visibles to data
2. Sample (  $h \sim p(h|v)$  )
3. Sample (  $v' \sim p(v|h)$  )

Purpose:

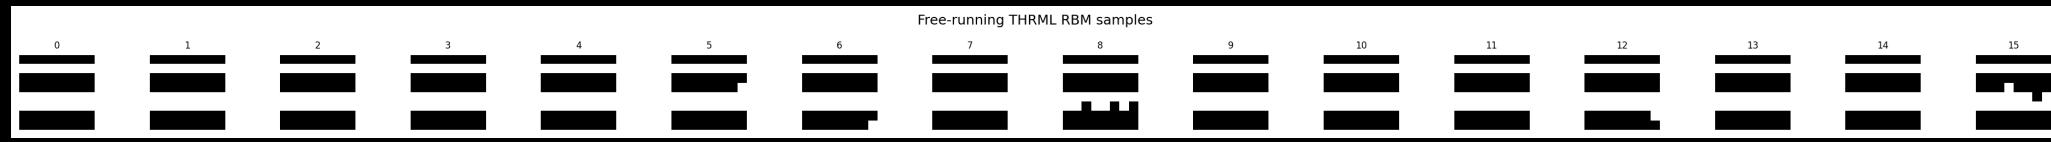
- Validate training
- Check encoder–decoder consistency
- Ensure stable attractor basins



# THML Free-Running Sampling

- Start from random spins
- Run full Gibbs chain (no clamping)
- THML explores the learned energy landscape
- Produces pure model samples

**Shows true generative behavior of the RBM.**



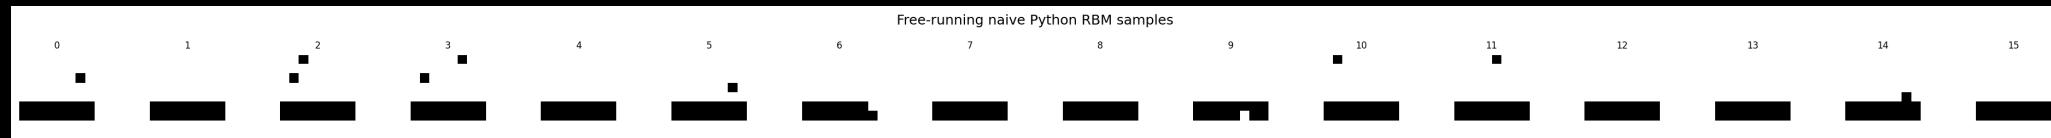
# Naive Python Gibbs Sampling (Baseline)

Why?

- Validate correctness of THML sampler
- Check if a hand-written sampler produces similar samples

Results:

- Python sampler produces Bars & Stripes
- Confirms correctness of learned energy landscape
- THML behaves consistently with standard RBM sampling



# Why Three Sampling Modes?

## 1. Clamped Reconstructions

- ✓ Measure training quality
- ✓ Checks encoder/decoder mapping
- ✓ Local consistency

## 2. THRML Free-Running Sampling

- ✓ Explore full Boltzmann distribution
- ✓ Checks generative behavior
- ✓ Shows learned modes

### 3. Naive Gibbs Sampling

- ✓ Baseline sanity check
- ✓ Confirms THMLL correctness
- ✓ Ensures consistent energy landscape

# What We Learned

## RBM Learning

- RBM trained via THRML achieves extremely low recon error
- Filters reflect dataset structure

## Energy Landscape

- THRML free-running samples match Bars & Stripes
- Energy minima match dataset modes

## Sampler Validation

- Python Gibbs sampler matches THRML outputs
- Confirms correctness of THRML's block sampler

# THRML as a Gibbs Engine

## Benefits

- Eliminates hand-written Gibbs loops
- Works for **arbitrary** factor graphs
- Supports multi-chain sampling
- Hardware-friendly (GPU/TPU acceleration)
- Clean integration with EBMs

## Takeaway

THRML can serve as a **general-purpose thermodynamic sampling engine** for probabilistic models.

# Conclusion

We successfully:

- Built a THRML-powered RBM
- Learned structured representations
- Generated high-quality samples
- Validated sampler correctness
- Demonstrated the feasibility of using THRML as a Gibbs engine

**A strong foundation for larger EBMs and future thermodynamic computing systems.**

# Thank You

Questions?

Contact: Alen Ribić / @alenribic, Alessio Toniolo / @\_alessiotoniolo